

Data Science for Business

Using pandas for Better (and Worse) Data Science

Asst. Prof. Teerapong Leelanupab (Ph.D.)
Faculty of Information Technology
King Mongkut's Institute of Technology Ladkrabang (KMITL)



Week 5.1

Using pandas for Better (and Worse) Data Science

Good Talk By [Kevin Markham](#) at [PyCon](#) Cleveland 2018

This tutorial is for intermediate pandas users

- Github: <https://github.com/justmarkham/pycon-2018-tutorial>
- Youtube playlist: <http://tiny.cc/t5i8az>

Datasets:

- police.csv is the Rhode Island dataset from the [Stanford Open Policing Project](#), made available under the [Open Data Commons Attribution License](#).
- ted.csv is the TED Talks dataset from [Kaggle Datasets](#), made available under the [CC BY-NC-SA 4.0 license](#).

Data Science for Business

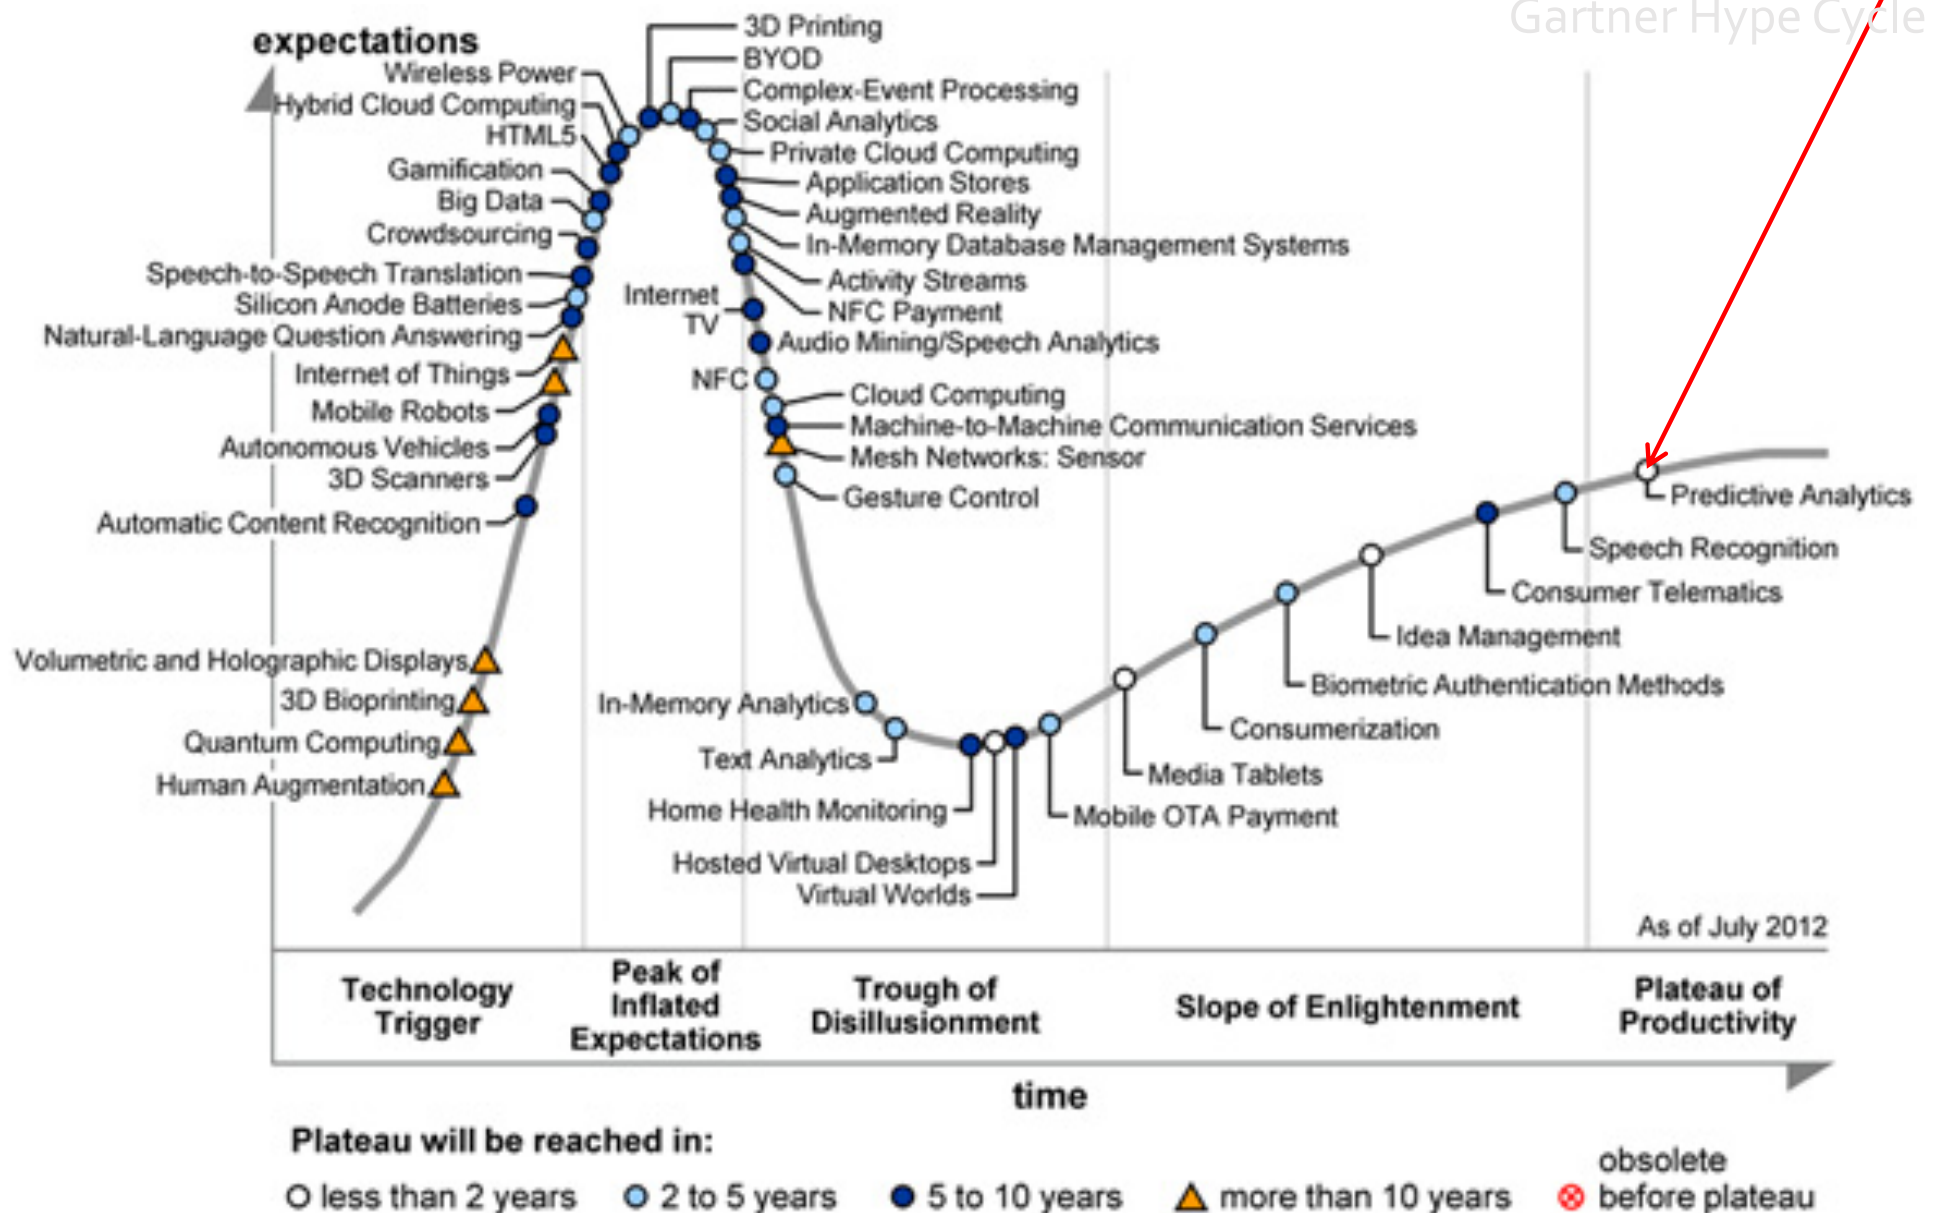
Predictive Modeling

Asst. Prof. Teerapong Leelanupab (Ph.D.)
Faculty of Information Technology
King Mongkut's Institute of Technology Ladkrabang (KMITL)



Week 5.2

Toward Predictive Analytics



VodafoneTelCo: Predicting Customer Churn



You just landed a great analytical job with VodafoneTelCo, one of the largest telecommunication firms in the US



They are having a major problem with customer retention in their wireless business

- In the mid-Atlantic region, 20% of cell phone customers leave when their contracts expire. Communications companies are now engaged in battles to attract each other's customers while retaining their own
- Marketing has already designed a special retention offer
- Your task is to devise a precise, step-by-step plan for how the data science team should use VodafoneTelCo vast data resources to solve the problem

VodafoneTelCo : Predicting Customer Churn

- What data you might use?
- How would they be used?
- How should VodafoneTelCo choose a set of customers to receive their offer in order to best reduce churn for a particular incentive budget?



Terminology

- Model:
 - A simplified representation of reality created to serve a purpose
- Predictive Model:
 - A formula for estimating the unknown value of interest: **the target**
 - The formula can be mathematical, logical statement (e.g., rule), etc.
- Prediction:
 - Estimate an unknown value (i.e. the target)
- Instance / example:
 - Represents a fact or a data point
 - Described by a set of **attributes** (fields, columns, variables, or features)

Terminology

- Model induction:
 - The creation of **models** from data
- Training data:
 - The input data for the induction algorithm

Terminology

Name	Balance	Age	Employed	Write-off
Mike	\$200,000	42	no	yes
Mary	\$35,000	33	yes	no
Claudio	\$115,000	40	no	no
Robert	\$29,000	23	yes	yes
Dora	\$72,000	31	no	no

This is one row (example).

Feature vector is: **<Claudio,115000,40,no>**

Class label (value of Target attribute) is **no**

What is a model?

A simplified representation of reality created for a specific purpose*

**based on some assumptions*


- *Examples: map, prototype, Black-Scholes model, etc.*
- *Data Mining example:*
 - "formula" for predicting probability of customer attrition at contract expiration*
 - "classification model" or "class-probability estimation model"*

Feature Types

- Numeric: anything that has some order
 - Numbers (that mean numbers)
 - Dates (that look like numbers ...)
 - Dimension of 1
- Categorical: stuff that does not have an order
 - Binary
 - Text
 - Dimension = number of possible values (-1)
- Food for thought: Names, Ratings, Standard Industrial Classification ([SIC](#))

Dimensionality of the data?

Attributes / Features



Name	Balance	Age	Default
Mike	\$123,000	30	Yes
Mary	\$51,100	40	Yes
Bill	\$68,000	55	No
Jim	\$74,000	46	No
Mark	\$23,000	47	Yes
Anne	\$100,000	49	No

- **Dimensionality of a dataset** is the sum of the dimensions of the features
 - The sum of the number of numeric features and ~ the number of values of categorical features

Supervised versus Unsupervised Methods

- “Do our customers naturally fall into different groups?”
 - No guarantee that the results are **meaningful** or will be useful for any particular purpose
- “Can we find groups of customers who have particularly high likelihoods of canceling their service soon after contracts expire?”
 - **A specific purpose**
 - Much more useful results (usually)
 - Different techniques
 - **Requires data on the target**
 - The individual’s label

Supervised Learning & Predictive Modeling

- Is there a specific, quantifiable **target** that we are interested in or trying to predict?
 - Think about the decision
- Do we have data on this target?
 - Do we have enough data on this target?
 - Need a min ~500 of each type of classification
 - Do we have relevant data prior to decision?
 - Think timing of decision and action
- The result of supervised data mining is a model that predicts some quantity
- A model can either be used to predict or to understand

Subclasses of Supervised Learning Technique

- Classification
 - Categorical target
 - Often binary
 - Includes “class probability estimation”
- Regression
 - Numeric target

Subclasses of Supervised Learning Technique

- “Will this customer purchase service S_1 if given incentive I_1 ?”
 - Binary Classification problem
 - Dichotomous target (the customer either purchases or does not)
- “Which service package (S_1 , S_2 , or none) will a customer likely purchase if given incentive I_1 ?”
 - Multiclass Classification problem
 - Three-valued target
- “How much will this customer use the service?”
 - Regression problem
 - Numeric target
 - Target variable: amount of usage per customer

Supervised Learning/Predictive Modeling

Key (part 1): is there a specific, quantifiable target that we are interested in or trying to predict?

Examples:

- What will the IBM stock price be tomorrow? (e.g., \$200)
What would you do if you could predict this?
- Will this prospect default her loan? (e.g. yes/no)
What would you do if you could predict this?
- Do my customers naturally fall into different groups?
[Unsupervised: no objective target stated.]
What would you do if you could predict this?

Supervised Learning/Predictive Modeling

*Key (part 2): do we have data on this target?
supervised learning requires both parts 1 & 2*

We don't need the exact data (e.g., whether the current customers will leave)
BUT we need data for the same or a related phenomenon (e.g., from
customers from last quarter)

→ we will use these data to build a model to predict the phenomenon of interest

- Think: Is the phenomenon of “who left the company” last quarter the same as the phenomenon of “who will leave the company” next quarter?
- Think: Who might buy this completely new product I have never sold before?

Supervised Learning/Predictive Modeling

*Key (part 3): the result of supervised learning is a **MODEL** that given data predicts some quantity*

- if (income <50K)
 then no Life Insurance
 else Life Insurance


[Result of Supervised: you can apply this rule to any customer and it gives you prediction]

What might a predictive model look like?

- There are different sorts of model. Here are just two examples:

- Tree/Rule: (a supervised segmentation)
- If (income > \$50K) & (age > 45) then LI = YES
- If ... then ...


- Numeric function:
- $P(LI) = f(x_1, x_2, \dots, x_k)$



- **Did not buy life insurance**

+ Bought life insurance

Classification tree

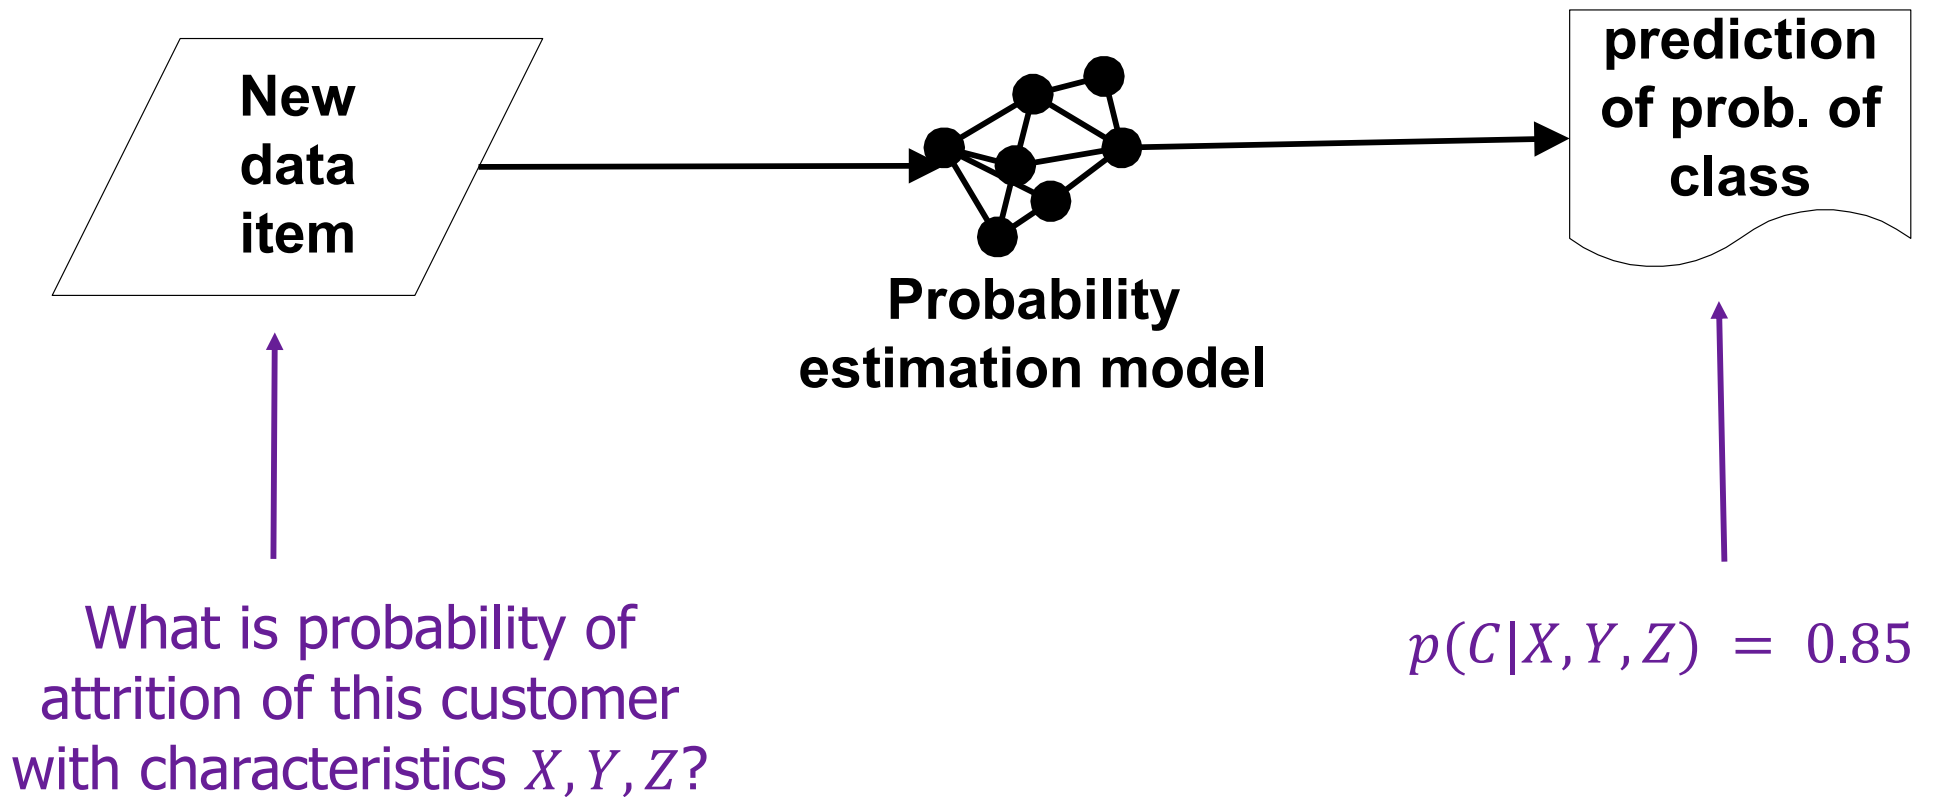


Supervised Data Mining/Predictive Modeling

*Key (part 4): a data-driven model can either be used to predict or to understand**

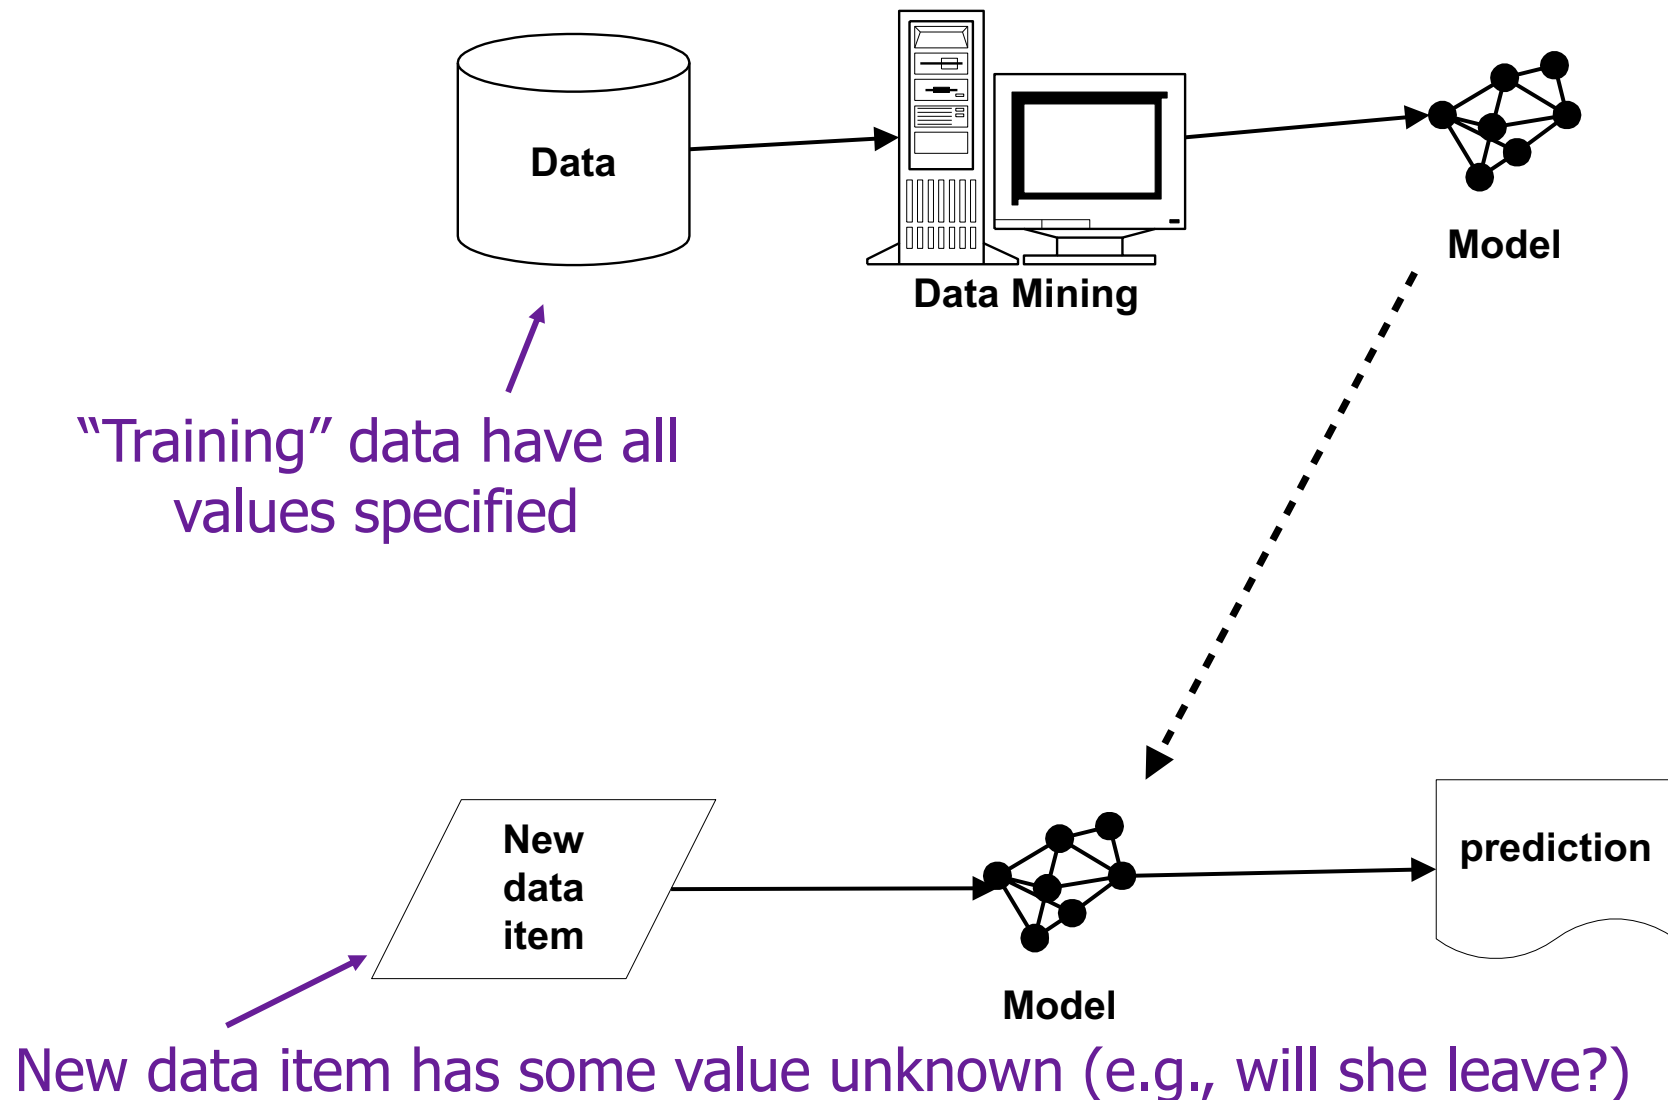
**Explanatory modeling can be quite complex. It turns out that you need to understand the fundamentals of predictive modeling first.*

Predictive Model in Use



Data Mining versus Use of the Model

"Supervised" modeling:



Classical Pitfalls in modeling setup

- The training data is NOT consistent with the use
- Bad sample
- Bad features

Sample: “Looking under the streetlight”

- Target Proxy
 - I do not see if a person after seeing an *ad* bought the book, so lets model clicks ...
- Sample Proxy
 - I want to run a campaign in Spain but only have data on US customers

https://en.wikipedia.org/wiki/Streetlight_effect

Sample: “Survivorship issues”

- Lending club wants to have a model to take over the screening (คัดกรอง) process that selects applications and deny those that are likely to default (ผิดชำระหนี้)
- Data of past loans and the outcomes are provided
- Bad:
 - Use only the data (without defaults on the loan), they currently have to predict default (credit default risk)

Recall Key 2: We don't need the exact data (e.g., whether the application of current customers who will repay their loans) BUT we need data for the same or a related phenomenon (e.g., from customers from last quarter)

→ *we will use these data to build a model to predict the phenomenon of interest*

Sample: Different Sources

- Things go really bad if the positive and negative are treated differently
- Looking for drivers of diabetes: how do you assemble the training data?
- Bad:
 - Go to a specialized hospital and get records from people treated for diabetes
 - Go somewhere else to get records for healthy people

Summary on bad habits

- You are missing all the applications that were turned down already
- The sick people came from a very artificial subset
- Your target is NOT really your target
- No way of telling how the model will perform
 - No way of testing either
- The training sample should be as similar as possible to the USE data

Data Science for Business

Modeling and Prediction in scikit-learn

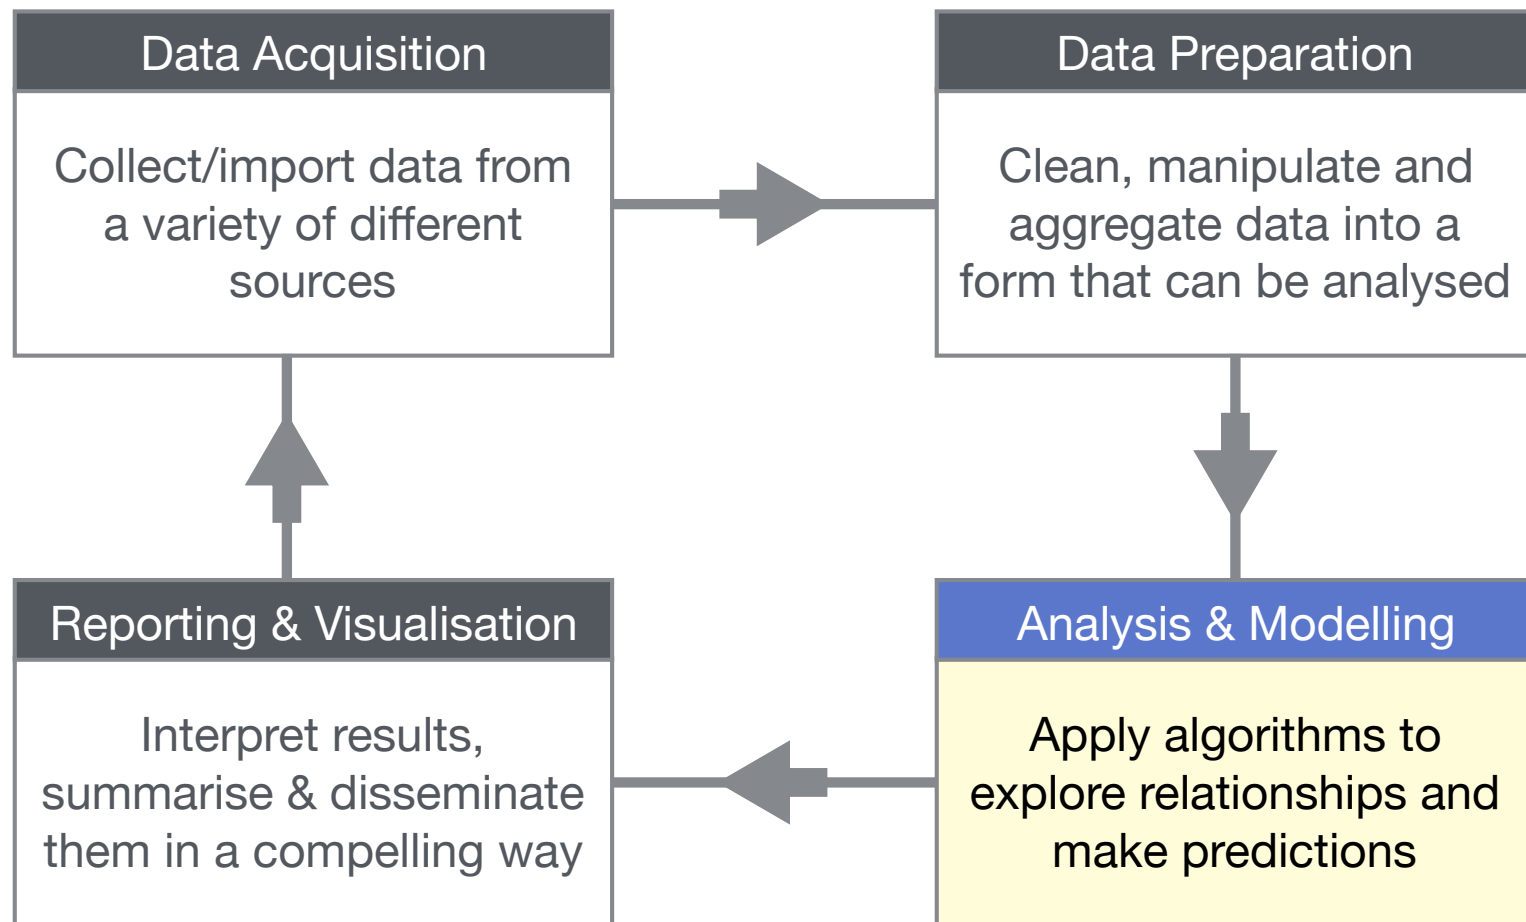
Asst. Prof. Teerapong Leelanupab (Ph.D.)
Faculty of Information Technology
King Mongkut's Institute of Technology Ladkrabang (KMITL)



Week 5.3

Reminder: Data Science Pipeline

- Recall the stages of the basic data science pipeline...
- Most complex component relates to data mining and modelling.



Modelling and Prediction Tasks

- **Predictive modelling** uses statistics to predict outcomes, based on historic data. Also referred to as **supervised machine learning**.
- **Examples:**
 - Spam filtering: predict if a new email is spam or non-spam, based on annotated examples of past spam / non-spam.
 - Car insurance: assign risk of accidents to policy holders and potential customers.
 - Healthcare: predict disease which a patient has, based on their symptoms.
 - Algorithmic trading: predictive models can be built for different assets like stocks, futures, currencies, etc, based on historic data and company information.

Supervised Learning

- Classification:**

Learn from a labelled training set to make a prediction to assign a new "unseen" example to one of a fixed number of classes.

Unseen
Input
Email

```
Received: from PACIFIC-CARRIER-ANNEX.MIT.EDU by pe9.MIT.EDU (5.61/4.7) id AA29041; Mon, 18 Jan 99
Received: from mail.arvidejaure.se by MIT.EDU with SMTP
id AA19666; Mon, 18 Jan 99 21:05:04 EST
Received: from updatasec ([ip245:update-se-quick.com [209.213.129.245]]) by mail (SMI-8.6/SMI-SV94)
id CAA10872; Tue, 19 Jan 1999 02:40:37 +0100
Date: Tue, 19 Jan 1999 02:40:37 +0100
From: richards@therage.com
Message-Id: <>
Subject: $1,000 for sending an e-mail
Apparently-To: <lots of omitted names>
Apparently-To: <postmaster@mit.edu>

We will give you $1,000 for sending an e-mail to your friends. AB Mailing,
Inc. is proud to announce the start of a new contest. Each day until
January, 31, 1999, one lucky Internet or AOL user who forwards our
advertisement to their friends will be randomly picked to receive $1,000!
You could be the winner!

Thank you for your time.

Our message:
"YOUR COMPUTER CAN MAKE MONEY WHILE YOU SLEEP"
Respected 10 Year Old Company Shows You How
YOU HAVE NEVER SEEN ANYTHING LIKE THIS!

http://www.abmailing.com
Phone: 800-226-2770
```

Model


Non-Spam

Spam

2 classes

- Regression:**

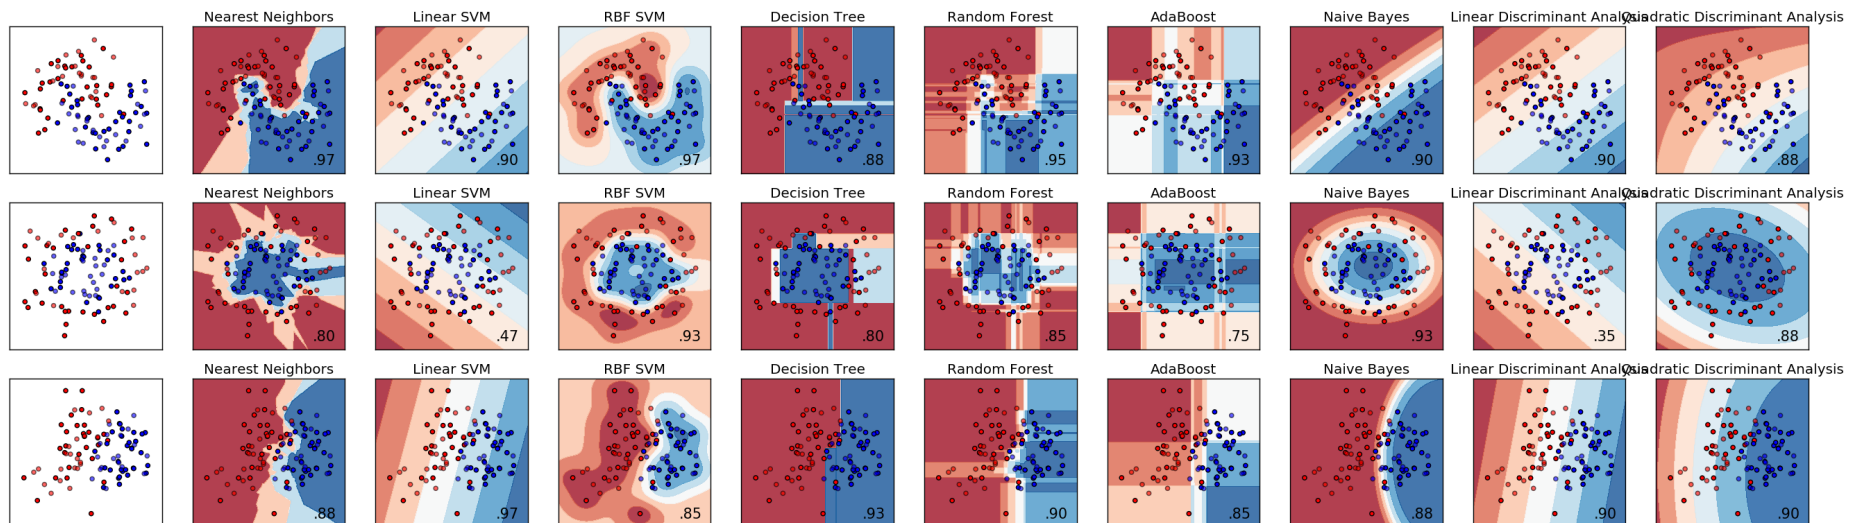
Learn from a labelled training set to decide the value of a continuous output variable (i.e. the output is a number).



Scikit-Learn Package

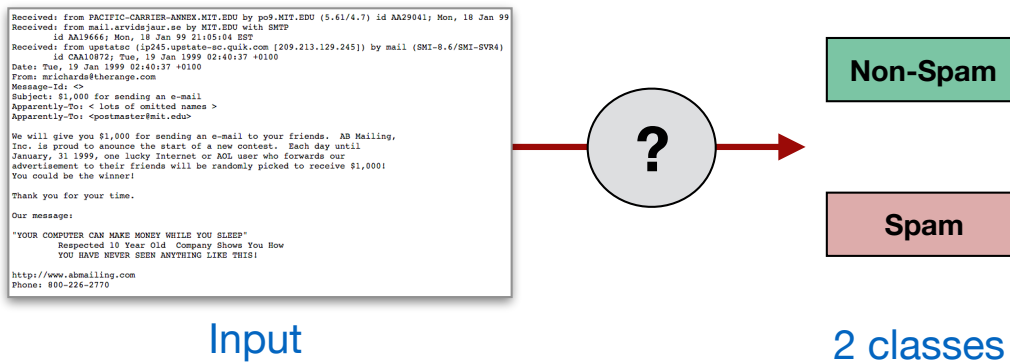
- Scikit-learn is a comprehensive open source Python package for machine learning and data analysis: <http://scikit-learn.org>
- Anaconda includes Scikit-learn as part of its free distribution.
- Scikit-learn algorithm inputs and outputs are usually represented as NumPy arrays, although we can also work with input data as Pandas Data Frames.

```
import sklearn
```



Types of Classification Tasks

- Binary Classification:**
Assign an input to one of two possible target class labels.



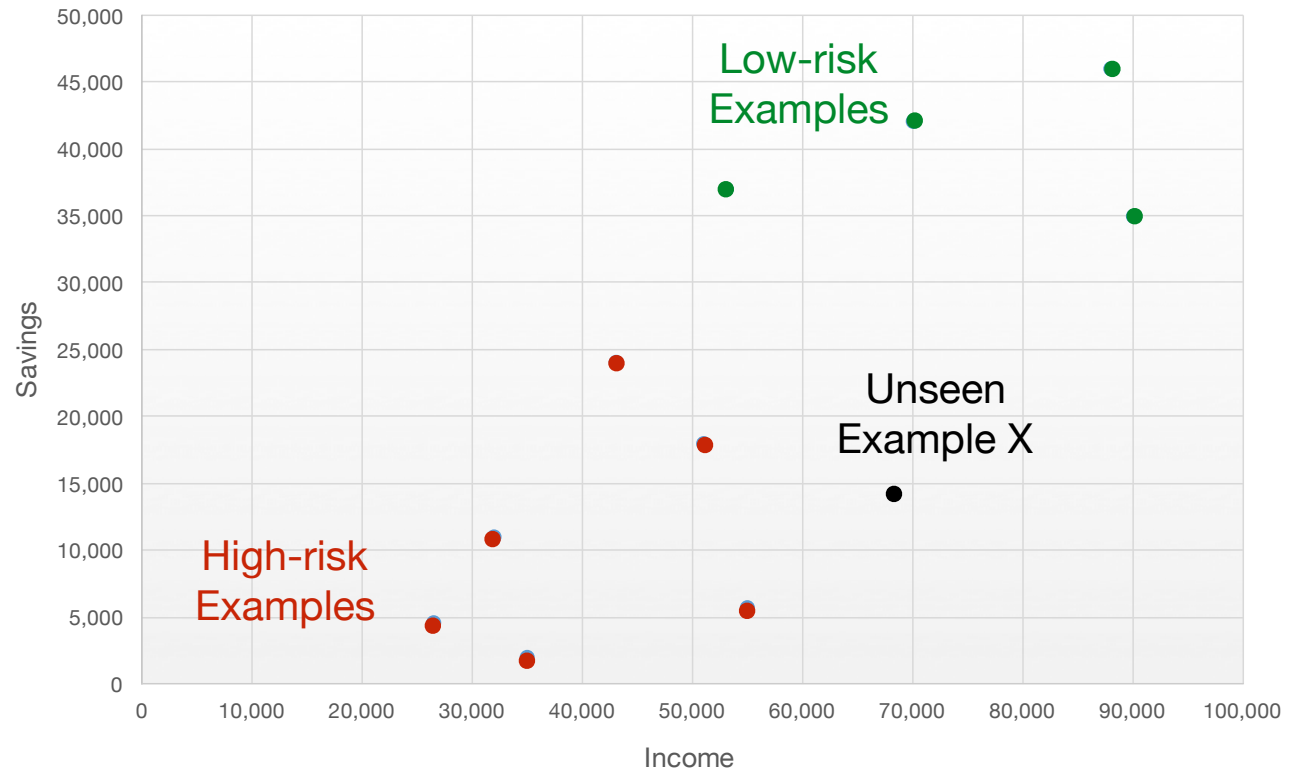
- Multiclass Classification:**
Assign an input to one of M different target class labels.



Typical Binary Classification Task

Example: Mortgage Approval

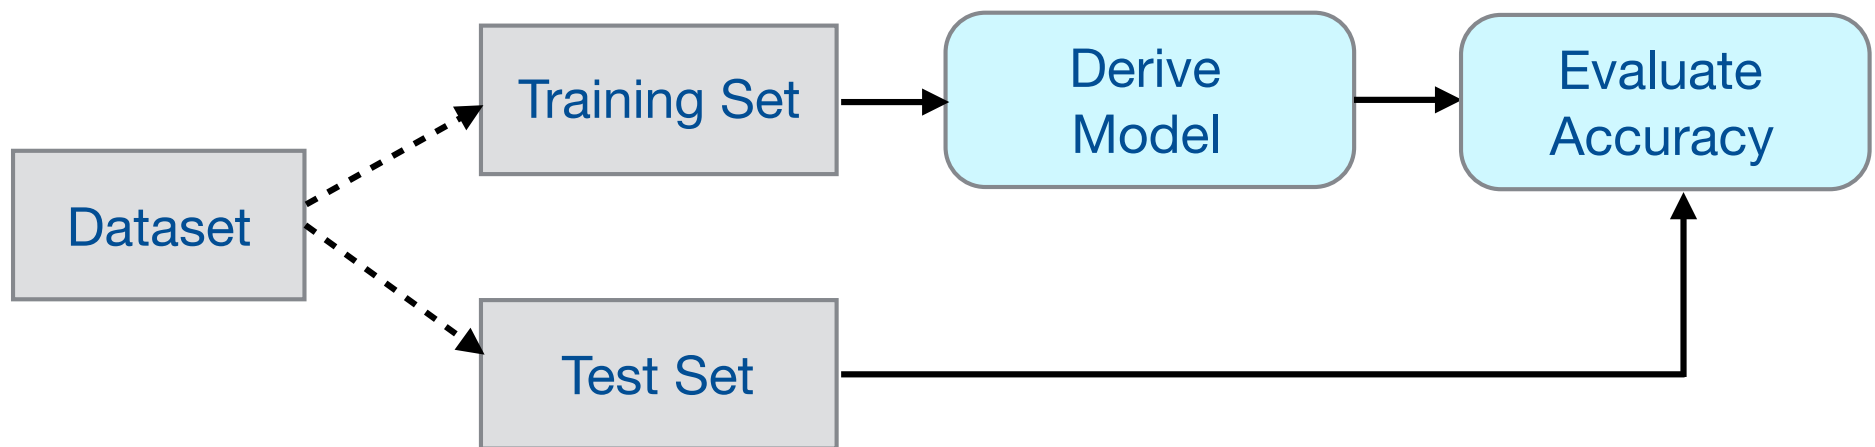
Manually classify mortgage applicants into two categories (**low-risk** & **high-risk**) based on savings and income data.



Q. Can we train an algorithm to learn to automatically classify a new application X as either **low-risk** or **high-risk**, based on the manually-classified examples of each type?

Training v Test Data

- Standard evaluation approach for classification tasks is to split the set of examples into a *training set* and the *test set*.
- **Training set:** Examples provided to the classifier to build a model of the data. Each has been manually assigned a class label.
- **Test set:** Examples held back from the classifier, which are used to evaluate the accuracy of the classifier. Test examples are completely separate from the training set, so we can evaluate how well the model built by the classifier will generalise to new input examples.



Classification Algorithms

- Many different learning algorithms exist for classification (e.g. k-nearest neighbour, decision tree, neural network, support vector machine).
- Problem dimensions will often determine which classification algorithm will be practically applicable, due to processing, memory, and storage requirements.
 1. Number of examples N .
 - Sometimes millions of input examples.
 2. Number of features (dimensions) D representing each example.
 - Often 10-1000, but sometimes far higher.
 3. Number of target classes M .
 - Often small (binary), but sometimes far higher.

Similarity/Distance-based Learning


Fundamental Strategy: “Best way to make predictions is to look at past examples and repeat the same process again”.

Feature space:

A D -dimensional coordinate space used to represent the input examples for a given problem, with one coordinate per descriptive feature.

Similarity measure:

Some function to measure how similar (or distant) two input examples are from one another are in the D -dimensional coordinate space.



2 features describing each example (agility & speed)

→ 2 coordinate dimensions for measuring similarity


Typically: Distance = $1/\text{Similarity}$ OR Distance = $1 - \text{Similarity}$

Similarity/Distance-based Learning

- For numeric data, the most common distance function used is the **Euclidean distance**.
- Calculated as square root of sum of squared differences for each feature f representing a pair of examples.

$$ED(\mathbf{p}, \mathbf{q}) = \sqrt{\sum_{f \in F} (q_f - p_f)^2}$$

Example	Speed	Agility
x4	3.25	8.25
x15	4.75	6.25




$$ED(x4, x15) = \sqrt{(3.25 - 4.75)^2 + (8.25 - 6.25)^2} = \sqrt{6.25} = 2.5$$

KNN Classifier


k -Nearest Neighbour Classifier: Simple but effective "lazy" classifier. Find most similar previous examples for which a decision has already been made (i.e. their **nearest neighbours** from the training set).

Example: For new unseen examples, look at the label of k nearest neighbours under both features (e.g. $k=3$ neighbours)



KNN Classifier

Majority voting: The decision on a label for a new query example is decided based on the “votes” of its k nearest neighbours, where the neighbours are selected to minimise the distance $d(q, x_i)$



Neighbour counts

- Yes = 2 votes
- No = 1 vote

➡ Majority says Yes!

KNN Classifier in Python

- Goal: Train a classifier to predict the species (class label) for a new numeric record describing a flower.
- Three possible species (class labels), so this is a multi-class problem: 'setosa', 'versicolor', or 'virginica'.
- Step 1: Fit a KNN classifier with K=3 neighbours on the full dataset - this is our training data, using the `fit()` function.

```
from sklearn.neighbors import KNeighborsClassifier
model = KNeighborsClassifier(n_neighbors=3)
model.fit(iris.data, iris.target)
```

- Step 2: Make a prediction for a new unseen input example using the trained KNN model, using the `predict()` function.

```
xinput = np.array([[3.0, 5.0, 4.1, 2.0]])
pred_class_number = model.predict(xinput)
print( iris.target_names[pred_class_number] )
```

```
['virginica']
```

KNN Classifier in Python

- We can also generate predictions for multiple input examples using a single call to `predict()` - i.e. apply classifier to a new test set.
- In general we will have 4 variables:
 1. `dataset_train`: The set of data on which to build the model.
 2. `target_train`: The true class labels for the training data.
 3. `dataset_test`: The set of data on which to test the model.
 4. `target_test`: The true class labels for the test data (if available).

```
model = KNeighborsClassifier(n_neighbors=3)
model.fit(dataset_train, target_train)
```

Fit a KNN classifier with K=3
neighbours on the training data

```
predicted = model.predict(dataset_test)
print(predicted)
```

Make a prediction for the test set

```
[ 1.  0.  0.  0.  1.  0.  0.  0.  0.  1.  0.  1. ]
```

A prediction is made for each
input in the test set

```
print(target_test)
```

```
[ 1.  0.  0.  0.  1.  0.  0.  1.  1.  0.  0.  1. ]
```

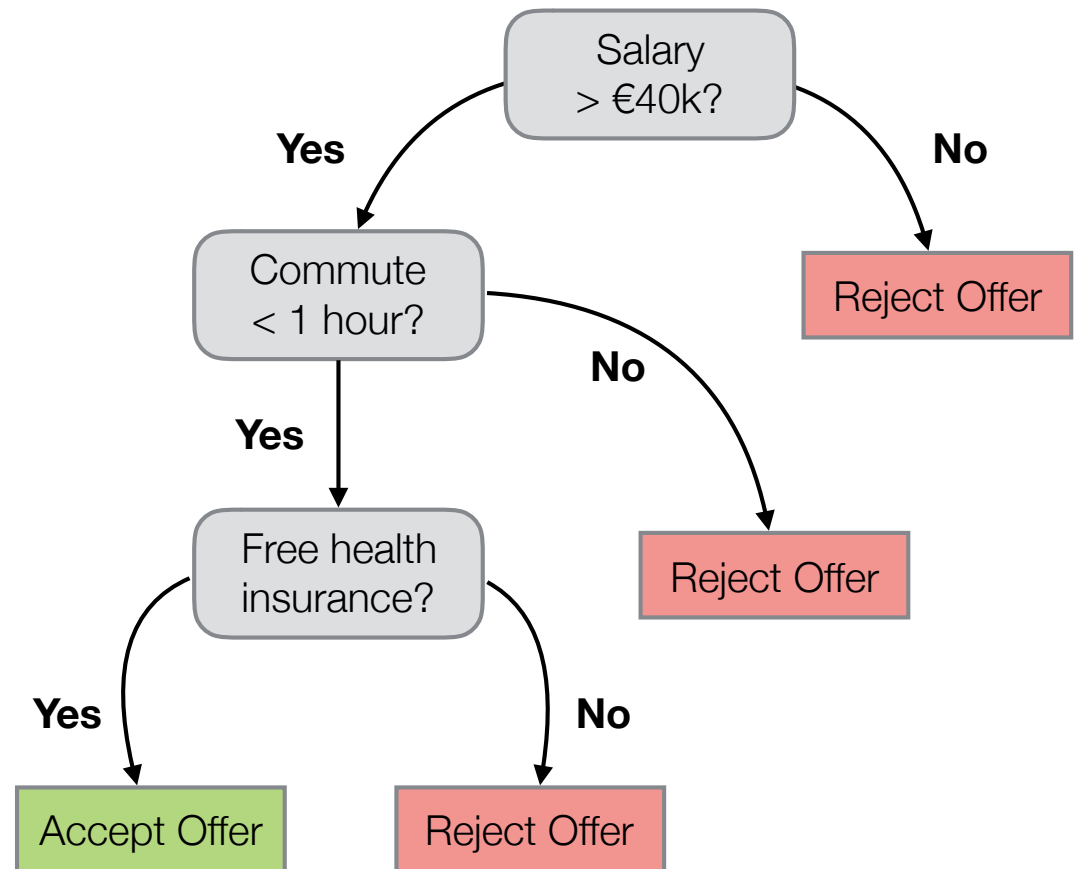
We can compare the predictions to
the true labels for the test data

Decision Tree Classifier

- Create a classification model to predict class labels by learning decision rules learned from training data.



Decision:
Should I accept the
job offer?



Decision Tree Classifier

- **Basic Idea:** Build a tree model that splits the training set into subsets in which all examples have the same class.
- Splitting is done using rules inferred from the training set.
- Each rule is based on a feature, and the split corresponds to the values it can take:
 - e.g. `insured = {true, false}`
 - e.g. `wind = {weak, strong}`
 - e.g. `income = {low, average, high}`
 - e.g. `height < 6ft, height ≥ 6ft`
- If necessary, each subset can be split again using another feature, and so on until all examples have the same class.
- Once the tree is built, we can use it to quickly classify new input examples - this is an eager learning strategy.

Decision Tree Example

Q. “Will a customer wait for a restaurant table?”

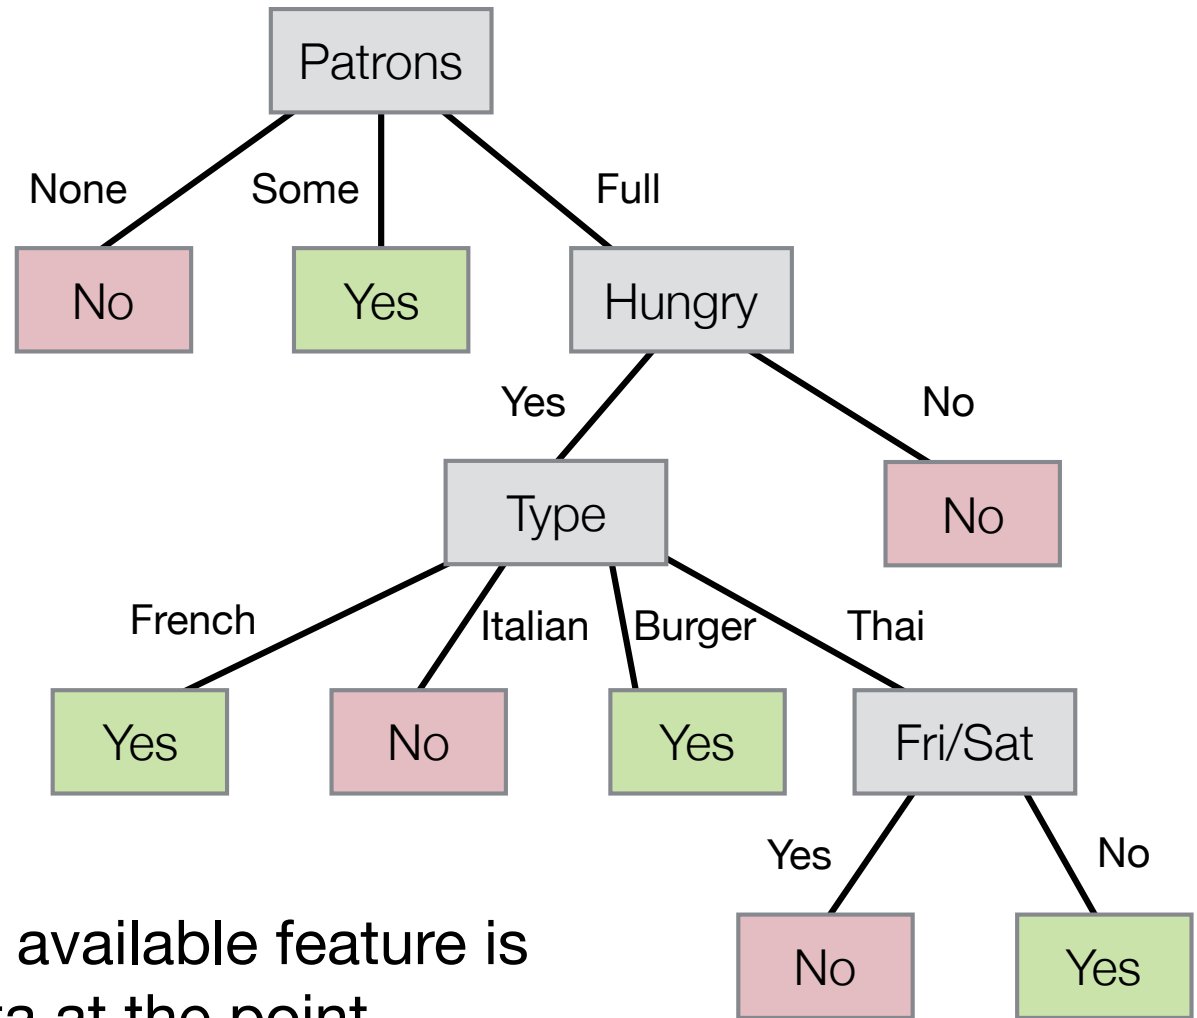
Russell & Norvig, Artificial Intelligence: A Modern Approach, Prentice Hall, 2009.

Binary classification task (WillWait = {Yes,No}), with examples described by 10 different features:

Example	Alternate	Bar	Fri/Sat	Hungry	Patrons	Price	Raining	Reservation	Type	WaitEst	WillWait?
1	Yes	No	No	Yes	Some	€€€	No	Yes	French	0-10	Yes
2	Yes	No	No	Yes	Full	€	No	No	Thai	30-60	No
3	No	Yes	No	No	Some	€	No	No	Burger	0-10	Yes
4	Yes	No	Yes	Yes	Full	€	No	No	Thai	10-30	Yes
5	Yes	No	Yes	No	Full	€€€	No	Yes	French	>60	No
6	No	Yes	No	Yes	Some	€€	Yes	Yes	Italian	0-10	Yes
7	No	Yes	No	No	None	€	Yes	No	Burger	0-10	No
8	No	No	No	Yes	Some	€€	Yes	Yes	Thai	0-10	Yes
9	No	Yes	Yes	No	Full	€	Yes	No	Burger	>60	No
10	Yes	Yes	Yes	Yes	Full	€€€	No	Yes	Italian	10-30	No
11	No	No	No	No	None	€	No	No	Thai	0-10	No
12	Yes	Yes	Yes	Yes	Full	€	No	No	Burger	30-60	Yes

Decision Tree Example

- A “good” decision tree will classify all examples correctly using as few tree nodes as possible.

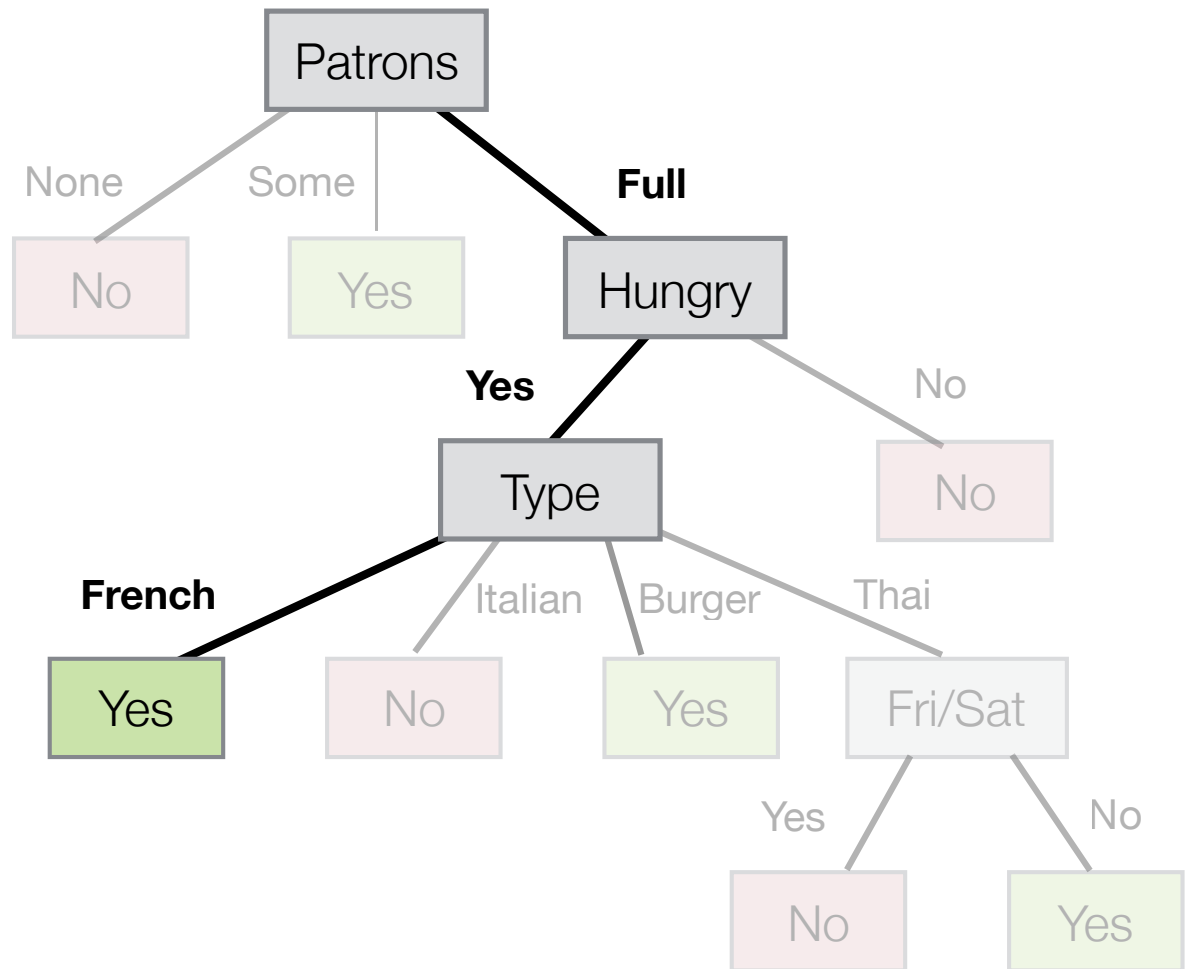


At each level, the best available feature is chosen to split the data at the point.

Decision Tree Example

- Once we have built a decision tree from the training set, we can use the rules in the tree to quickly classify new input examples.

Feature	Query <i>x1</i>
<i>Alternate</i>	Yes
<i>Bar</i>	No
<i>Fri/Sat</i>	Yes
<i>Hungry</i>	Yes
<i>Patrons</i>	Full
<i>Price</i>	€€
<i>Raining</i>	No
<i>Reservation</i>	No
<i>Type</i>	French
<i>WaitEstimate</i>	10-30



⇒ Based on tree, output for *x1* is "Yes"

Decision Tree Classifier in Python

- **Example:** Apply a Decision Tree to the *Iris* dataset.
- Step 1: Apply a Decision Tree classifier on the full dataset, again using the `fit()` function to create the model.

```
from sklearn.datasets import load_iris
from sklearn.tree import DecisionTreeClassifier
iris = load_iris()
model = DecisionTreeClassifier()
model.fit(iris.data, iris.target)
```

- Step 2: Make a prediction for a new unseen input example using the trained tree model, using the `predict()` function.

```
xinput = np.array([[2.9, 5.1, 4.1, 1.8]])
pred_class_number = model.predict(xinput)
print( iris.target_names[pred_class_number] )
```

```
['versicolor']
```

Alternative Classifiers

- **Naive Bayes:** Simple classifier based on counts. Requires less training data, assumes all features are independent.
- **Support Vector Machines:** SVMs use a kernel function to accurately separate different classes. Difficult to interpret, can require parameter tuning.
- **Logistic Regression:** Probabilistic model, works well in an online setting.

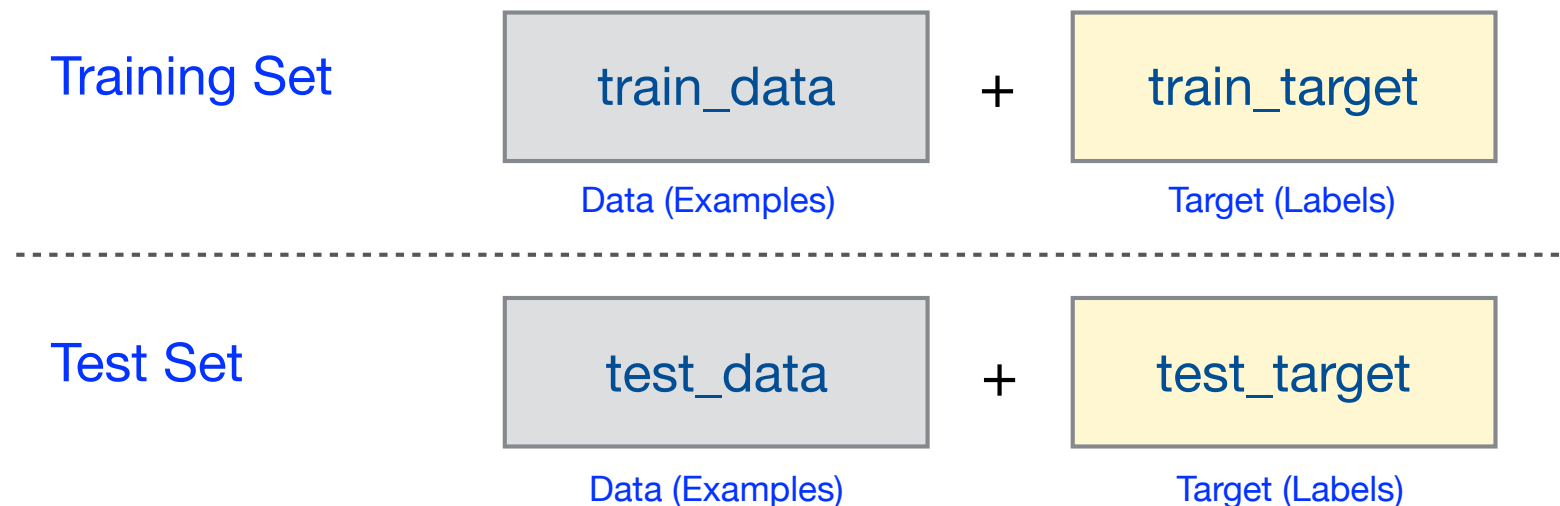
```
from sklearn.naive_bayes import MultinomialNB
model = MultinomialNB()
model.fit(data_train, target_train)
predicted = model.predict(data_test)
```

```
from sklearn.svm import SVC
model = SVC()
model.fit(data_train, target_train)
predicted = model.predict(data_test)
```

```
from sklearn import linear_model
model = linear_model.LogisticRegression()
model.fit(data_train, target_train)
predicted = model.predict(data_test)
```

Evaluating Classifiers

- Reminder - Standard evaluation approach for classification tasks is to split the set of examples into two sets:
 1. **Training set:** Examples provided to the classifier to build a model of the data. Each has been assigned a class label.
 2. **Test set:** Separate set of examples which are used to evaluate the accuracy of the classifier.
- **Hold-out strategy:** Randomly hold back some examples (e.g. 20%) for evaluation. In Scikit-learn, this split involves 4 variables:



Evaluation Measures

When making predictions, we need some kind of measure to capture how often the model makes correct or incorrect predictions, and how severe the mistakes are.











Accuracy:

Simplest measure. Fraction of correct predictions made by the classifier.

$$ACC = \frac{\# \text{ correct predictions}}{\text{total predictions}}$$

Example: Predictions made for test set of 10 emails

$$ACC = \frac{7}{10} = 0.7$$

Email	Label	Prediction	Correct?
1	spam	non-spam	
2	spam	spam	
3	non-spam	non-spam	
4	spam	spam	
5	non-spam	spam	
6	non-spam	non-spam	
7	spam	spam	
8	non-spam	spam	
9	non-spam	non-spam	
10	spam	spam	

Measuring Accuracy in Python

- The `sklearn.metrics` package provides evaluation functionality.
- To demonstrate, we will create an artificial dataset, where the examples have 2 binary class labels $[-1,1]$:

```
from sklearn.datasets import make_hastie_10_2
data, target = make_hastie_10_2(100)
```

100 examples in total, both the data and the target labels

```
from sklearn.model_selection import train_test_split
data_train, data_test, target_train, target_test =
train_test_split(data, target, test_size=0.2)
```

Randomly split the data into training and test sets. Test set size is 20% (0.2)

```
from sklearn.neighbors import KNeighborsClassifier
model = KNeighborsClassifier(n_neighbors=3)
model.fit(data_train, target_train)
```

Apply KNN classifier to training data with target labels

```
predicted = model.predict(data_test)
```

Make predictions for test data

```
from sklearn.metrics import accuracy_score
accuracy_score(target_test, predicted)
```

```
0.65
```

Calculate accuracy score for the predictions, based on actual target labels for test data